

CS3243 Notes

Contents

1	Lecture 1: Introduction	4
1.1	Intelligent Agents	4
1.2	Rationality	4
1.3	Task Environment: PEAS	4
1.4	Properties of Task Environments	4
1.5	Building an Agent	5
2	Lecture 2: Uninformed Search	6
2.1	(★) Search Problem Formulation	6
2.2	Searching for Solutions	6
2.3	Search Strategies	7
2.4	Breadth-First Search (BFS)	7
2.5	Uniform-Cost Search (UCS)	8
2.6	Depth-First Search (DFS)	8
2.7	Depth-Limited Search (DLS)	8
2.8	Iterative Deepening Search (IDS)	8
2.9	Choosing a Search Strategy	9
2.10	Search Tracing Problems	9
3	Lecture 3: Informed Search	10
3.1	Best-First Search	10
3.2	Greedy Best-First Search (special case of Best-FS)	10
3.3	A* Search (special case of Best-FS)	10
3.4	Heuristic Design	11
3.5	Local Search	11
4	Lecture 4: Adversarial Search	13
4.1	Adversarial Search Problems (Games)	13
4.2	Strategies	13
4.3	Optimal Decisions (Minimax)	14
4.4	α - β Pruning	14
4.5	Imperfect, Real-Time Solutions	17
5	Lecture 5: Constraint Satisfaction Problems	18
5.1	CSP Formulation	18
5.2	CSP Variants	19
5.3	Constraint Variants	19
5.4	CSP Search	21
5.5	Backtracking Search Algorithm	21
5.6	Backtracking Heuristics: Variable and Value Ordering	22
5.7	Inference: INFERENCE	23
6	Lecture 6: Project Details	25
6.1	Reinforcement Learning	25
6.2	Supervised vs Unsupervised Learning	25
6.3	Evaluating Classification Models	25
6.4	Algorithm Selection	25
7	Lecture 7: Constraint Satisfaction Problems II	27
7.1	Inference in CSPs: Arc Consistency and AC-3	27
8	Logical Agents	29

8.1	Knowledge-based Agents	29
8.2	Example: Wumpus World	30
8.3	Logic	31
8.4	Validity and Satisfiability	36
8.5	Applying Inference Rules	36
8.6	Resolution (for CNF)	36
8.7	KB and Horn Clauses	37
8.8	Forward Chaining	37
8.9	Backward Chaining	39
8.10	Forward vs Backward Chaining	39
9	Uncertainty	40
9.1	Probability Basics	40
9.2	Bayesian Inference	40
10	Bayesian Networks	41
10.1	Examples	41
10.2	Inference in Bayesian Networks	44
10.3	Algorithm for Constructing Bayesian Network	44
10.4	Markov Blanket	45
10.5	d-Separation	45

1 Lecture 1: Introduction

1.1 Intelligent Agents

Agents interact with their environment

- Sensors take in percepts
- Actuators perform actions
- Agent function maps *percept histories* to *actions*: $f : P^* \rightarrow A$

1.2 Rationality

Rational if selected actions are:

- Based on evidence (prior knowledge/percept sequence)
- Maximise performance measure

Performance measure: defining and measuring 'performance' is difficult

- Task specificity: easier to define 'performance' for a narrower than more general task

Can be rational to explore (perform actions that gather information)

Agent is *autonomous* if behaviour is determined by its own experience

1.3 Task Environment: PEAS

PEAS: Performance measure, Environment, Actuators, Sensors

E.g. Automated Taxi

- Performance measure: safe, fast, comfort, revenue
- Environment: roads, traffic, pedestrians
- Actuators: steering wheel, accelerator, brake
- Sensors: sonar, speedometer, gps, engine sensors

1.4 Properties of Task Environments

- Observability: fully or partially observable? (e.g. fog of war)
- Deterministic vs. stochastic: are there random elements?
 - Still deterministic if random elements do not affect the transition function
 - Not deterministic if some elements are unobservable to player
- Episodic vs. sequential
 - Episodic: choice of current action does not depend on actions in past episodes
 - Sequential: need to consider previous actions too (e.g. chess); current action affects future actions
 - *Order* is important in sequential, not in episodic
- Static vs. dynamic: is environment changing as agent deliberates?
- Discrete vs. continuous: finite/infinite number of distinct states/percepts/actions
 - We prefer solving discrete problems
- Single vs. multi agent

1.5 Building an Agent

Lookup table agent

- For each possible percept, write its optimal action
- Problem: huge table with many many possible percepts
- Problem: no autonomy, hard to change on-the-fly if action is wrong. Unmaintainable and rigid

Types of agents (in increasing complexity):

1. Simple reflex agent: passive, only acts when it observes a percept
2. Model-based reflex agent: passive, has state/internal model of the world
3. Goal-based agent: not just passive and based on percept; has goals and acts to achieve them
4. Utility-based agent: has utility function, acts to maximise it

State is updated based on percept, current state, most recent action, model of the world

(*) Utility function is *internal*, performance measure is *external* and used to assess agent

Learning agent

- Critic + learner => adapt based on performance standard

Exploration vs. exploitation: a classic trade-off the agent must make

- Exploration: get more knowledge to improve future gains
- Exploitation: make use of knowledge to maximise current gains

2 Lecture 2: Uninformed Search

Problem-solving agent: one kind of goal-based agent

Environment: fully observable, deterministic, discrete

Uninformed search: no additional knowledge incorporated

2.1 (★) Search Problem Formulation

- State: including initial state
 - Abstract ONLY the relevant information, and nothing else
 - Everything in the state should be a variable that can change, no constants
- Actions: $ACTIONS(s)$ gives set of all valid actions that can be executed in state s
 - Define it for every possible state s
- Transition model: $RESULT(s, a)$ gives new state s' upon doing action a in state s
 - Define it for every possible state s and its valid action a
- Goal test: test if a state s is the goal state
 - E.g. $IsCheckmate(s)$ or $IsSolved(s)$
- Path cost: path cost is additive sum of step costs
 - Step cost $c(s, a, s')$ — e.g. 1 per action taken

2.2 Searching for Solutions

Solution: sequence of actions leading from initial to goal state

Example: route planning

- Reduce map down to nodes with edges between them of certain weights

Example: 8-puzzle

- State: an arrangement of numbers in 3x3 grid, represented as matrix/array
- Actions: moving one filled square to a blank adjacent square
- Transition model: [depends on representation] — function that takes in state + action \Rightarrow new state
- Goal test: whether each cell matches the goal state, one-for-one
- Cost function: uniform cost of 1 for each action

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

State vs node

- State: represents physical configuration
- Node: data structure constituting part of search tree: includes state, parent node, action, path cost $g(n)$
- Two different nodes can contain the same world state

2.3 Search Strategies

Which order should we expand the nodes in?

Evaluation criteria

- Completeness: always find a solution if it exists
- Optimality: finds a least-cost solution
- Time complexity: number nodes generated
- Space complexity: max number of nodes in memory

Problem parameters

- b : maximum # of successors for each node — branching factor
- d : depth of *shallowest* goal node
- m : maximum depth of search tree

2.4 Breadth-First Search (BFS)

Frontier: Queue

Properties of BFS

- Complete: yes, as long as b is finite
- Optimal: no, unless uniform step cost, or uniform across each level
- Time: $O(b^d) = O(b) + O(b^2) + \dots + O(b^d)$
- Space: $O(b^d)$ (max size of frontier)

Applies goal test when pushing to frontier: reduces time and space complexity from $O(b^{d+1})$ to $O(b^d)$

2.5 Uniform-Cost Search (UCS)

Frontier: Priority queue, by path cost

- Idea: explore unexpanded node with *least-path-cost* (equivalent to BFS if all step costs are equal)

Properties of UCS

- Complete: yes, if all step costs are $\geq \epsilon$
 - If not, ever-decreasing step costs may get you stuck infinitely on a suboptimal path
 - Still yes even if b or d is infinite, or search space is infinite
- Optimal: yes (when it is complete)
- Time: $O(b^{1+\lfloor \frac{C^*}{\epsilon} \rfloor})$ where C^* is the optimal cost
 - Reach nodes at distance $0, \epsilon, 2\epsilon, \dots, \lfloor \frac{C^*}{\epsilon} \rfloor \epsilon$ of goal \Rightarrow total $\lfloor \frac{C^*}{\epsilon} \rfloor + 1$ steps
- Space: $O(b^{1+\lfloor \frac{C^*}{\epsilon} \rfloor})$

2.6 Depth-First Search (DFS)

Frontier: Stack

Properties of DFS

- Complete: yes, as long as depth is finite
- Optimal: no
- Time: $O(b^m)$
- Space: $O(bm)$ (can be $O(m)$ — at each level, just keep track of self and parent)

2.7 Depth-Limited Search (DLS)

Idea: run DFS with depth limit ℓ

- Only works if we know the goal is within ℓ steps
- Time: $O(b^\ell)$
- Space: $O(b\ell)$ (can be $O(\ell)$)

2.8 Iterative Deepening Search (IDS)

Idea: keep performing DLSs with increasing depth limit, until goal node is found

- Better if state space is large and depth of solution is unknown
- It can be wasteful with repeated effort
- But overhead is not that large (e.g. $b = 10, d = 5$ — 11%)

Properties of IDS

- Complete: yes, if b is finite
- Optimal: no, unless step cost is uniform
- Time: $O(b^d)$

- Space: $O(bd)$ (can be $O(d)$)

Property	BFS	UCS	DFS	DLS	IDS
Complete	Yes ¹	Yes ²	No	No	Yes ¹
Optimal	No ³	Yes	No	No	No ³
Time	$O(b^d)$	$O\left(b^{1+\left\lceil\frac{C^*}{\epsilon}\right\rceil}\right)$	$O(b^m)$	$O(b^\ell)$	$O(b^d)$
Space	$O(b^d)$	$O\left(b^{1+\left\lceil\frac{C^*}{\epsilon}\right\rceil}\right)$	$O(bm)$	$O(b\ell)$	$O(bd)$

1. Complete if b is finite
2. Complete b is finite and step cost $\geq \epsilon$
3. Optimal if step costs are identical

2.9 Choosing a Search Strategy

Depends on the problem

- Depth: finite/infinite?
- Solution depth: known/unknown?
- Repeated states
- Step costs: identical/different?
- Completeness and optimality – are they needed?
- Resource constraints (time/space)?

2.10 Search Tracing Problems

Tree-Search

```

Frontier
-----
S(0)
A(1) B(5) C(15)
S(2) B(5) G(11) C(15)
...

```

Graph-Search

```

Frontier      Explored
-----
S(0)
A(1) B(5) C(15)  S
B(5) G(11) C(15) S, A
G(10) C(15)      S, A, B

```

3 Lecture 3: Informed Search

Informed search: exploits problem-specific knowledge, uses *heuristics* to guide search

(AIMA Chapter 3.5.1-2, 3.6.1-...)

3.1 Best-First Search

Idea: use an *evaluation function* $f(n)$ for each node n

- Measures *cost estimate*
- Expand node with the lowest estimated cost first

Implementation: priority queue, ordered by non-decreasing cost f

3.2 Greedy Best-First Search (special case of Best-FS)

Evaluation function: $f(n) = h(n)$

- $h(n)$: cost estimate from n to goal (heuristic)
- Idea: expand the node that appears the closest to goal

Properties

- Complete: yes, if b is finite
- Optimal: no
- Time: $O(b^m)$, but if heuristic is good can reduce complexity substantially
- Space: $O(b^m)$ (max size of frontier)

3.3 A* Search (special case of Best-FS)

Idea: avoid expanding paths that are already expensive

- Expand the path that appears the cheapest

NOTE: remember we use a *priority queue* on $f(n) = g(n) + h(n)$; pick the smallest one

Evaluation function: $f(n) = g(n) + h(n)$

- $g(n)$: cost of reaching n from start node, under the current path (not necessarily the smallest among all paths!)
- $h(n)$: cost estimate from n to goal (heuristic)
- $f(n)$: estimated cost of cheapest path *through* n to goal

Properties

- Complete: yes, if there is finite number of nodes and $f(n) \leq f(G)$
- Optimal: yes, if you have an admissible/consistent heuristic
- Time (no great detail): $O(b^{h^*(s_0)-h(s_0)})$ where $h^*(s_0)$ is actual cost from root to goal
- Space: $O(b^m)$ (max size of frontier)

3.4 Heuristic Design

3.4.1 Admissibility

Admissible heuristics

- $h(n)$ is *admissible* if it never overestimates the cost to reach goal
- Definition: $\forall n, h(n) \leq h^*(n)$, where $h^*(n)$ is the true cost from n to goal state

Theorem: if $h(n)$ is admissible, then A* using TREE-SEARCH is optimal

- (Proof: see lecture 3 slide 22)

3.4.2 Consistency

Consistent heuristic:

- $h(n)$ is *consistent* if it means that $f(n)$ is non-decreasing along any path (triangle inequality)
- Definition: $h(n) \leq d(n, n') + h(n')$, where n' is a successor of n
- Lemma: if h is consistent, then $f(n') \geq f(n)$
- (???)

Theorem: if $h(n)$ is consistent, then A* using GRAPH-SEARCH is optimal

- Claim: when A* selects a node n for expansion, the shortest path to n has been found
- (Proof: see lecture 3 slide 26)

3.4.3 Admissibility & Consistency

All consistent heuristics are admissible, but not the other way round.

Example: 8-puzzle

- Heuristic 1: number of misplaced tiles
- Heuristic 2: total Manhattan distance

3.4.4 Dominance

h_2 *dominates* h_1 if $h_2(n) \geq h_1(n)$ for all n , where both heuristics are admissible

- Dominating heuristics are better: incur lower search costs under A*

3.4.5 Deriving Admissible Heuristics

Common exam question: given a problem, derive an admissible heuristic

Solution: *relax* the problem — then it'll only be 'easier' to reach the goal. Heuristic that uses this relaxed problem can NEVER over-estimate goal

3.5 Local Search

Path to the goal is irrelevant; we only want to reach the goal state

Local search algorithms: maintain single "current best" state, and try to improve it

Advantages

- Very little/constant memory
- Find reasonable solutions in large state space

3.5.1 Hill-Climbing Algorithm

HILL-CLIMBING

- $\text{current} \leftarrow \text{initial state}$
- while True:
 - $\text{neighbour} \leftarrow \text{best successor of current}$
 - if neighbour's value \leq current's value: return current
 - $\text{current} \leftarrow \text{neighbour}$

Problem: depending on initial state, can get stuck in local maxima (or minima)

Solution: try random restarts or sideways moves

4 Lecture 4: Adversarial Search

4.1 Adversarial Search Problems (Games)

Game: agent vs. agent(s)

- Unlike a *search problem*, which is agent vs. environment
- There are other utility-maximising agents
- Solution is a strategy that specifies a move for every possible opponent response

Zero-sum game: agent utilities sum to zero

- Completely adversarial game

Two-player zero-sum game

- *MAX* player: wants to maximise value
- *MIN* player: wants to minimise value

Problem formulation

- Initial state s_0
- States s
- (★ NEW) *Player* $\text{PLAYER}(s)$: defines which player has the move in state s
- Actions $\text{ACTIONS}(s)$: returns set of legal moves in state s
- Transition model $\text{RESULT}(s, a)$: returns state that results from move a in state s
- Terminal test $\text{TERMINAL}(s)$: check whether the game has ended
- (★ NEW) Utility function $\text{UTILITY}(s, p)$: final numeric value for game with terminal state s for player p

For now, we assume 2-player, deterministic, turn-taking

4.2 Strategies


Strategy

- Strategy s for player i : for every node of the tree that the player can make a move in, specify what player will do
- Need to define strategy in states that will never be reached (I think this means instead that it needs to be defined for all possible states)

Winning strategy

- Winning: s_1^* for player 1 is *winning* — if for any strategy s_2 by player 2, game ends with player 1 as the winner
- Non-losing: t_1^* for player 1 is *non-losing* — if for any strategy s_2 by player 2, game ends with EITHER player 1 as the winner or tie

4.3 Optimal Decisions (Minimax)



MINIMAX(s)

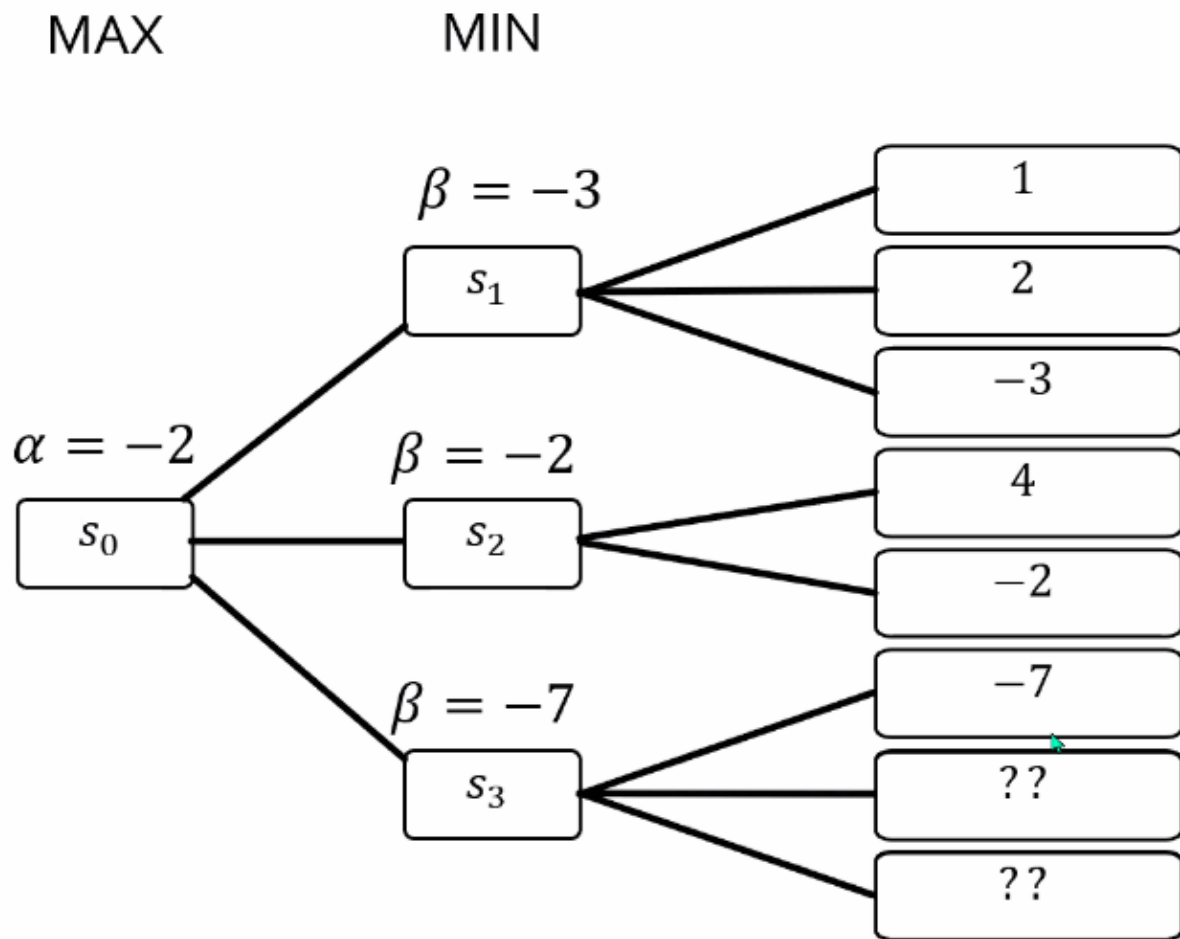
- UTILITY(s) if TERMINALTEST(s)
- $\max_{a \in \text{ACTIONS}(s)} \text{MINIMAX}(\text{RESULT}(s, a))$ if PLAYER(s) = MAX
- $\min_{a \in \text{ACTIONS}(s)} \text{MINIMAX}(\text{RESULT}(s, a))$ if PLAYER(s) = MIN

Properties


- Complete: yes, if game tree is finite
- Optimal: yes
- Time: $O(b^m)$ (similar to DFS)
- Space: $O(bm)$ (similar to DFS)

4.4 α - β Pruning

- α : largest value so far for MAX
- β : smallest value so far for MIN



Example above: in the bottom branch, $\beta = -7$, but $\alpha = -2 > \beta$. So no need to explore the remaining branches



α - β pruning

- MAX node n : $\alpha(n)$ = highest observed value found on path from n . Initially $\alpha(n) = -\infty$
- MIN node n : $\beta(n)$ = lowest observed value found on path from n . Initially $\alpha(n) = -\infty$
- (★) Given MIN node n , stop searching below n if there is some MAX ancestor i of n with $\alpha(i) \geq \beta(n)$
- (★) Given MAX node n , stop searching below n if there is some MIN ancestor i of n with $\beta(i) \leq \alpha(n)$

Analysis of α - β pruning

- "Perfect" ordering: time complexity = $O(b^{\frac{m}{2}})$ — can search twice as deep!
- Random ordering: time complexity = $O(b^{\frac{3}{4}m})$ for $b < 1000$

Summary

- Initially, $\alpha(n) = -\infty$, $\beta(n) = +\infty$
- $\alpha(n)$ is MAX along search path containing n
- $\beta(n)$ is MIN along search path containing n
- If a MIN node has value $v \leq \alpha(n)$, no need to explore further

- If a MAX node has value $v \geq \beta(n)$, no need to explore further

4.5 Imperfect, Real-Time Solutions

Time limit

- How to deal with super large search trees? \Rightarrow Limit maximum depth of tree
- Evaluation function: estimated expected utility of state (similar to heuristic)
- Cutoff test: e.g. depth limit

Cutting-Off Search: similar to Depth-Limited Search (DLS)

- Previously: $\text{MINIMAX}(s) = \text{UTILITY}(s)$ if $\text{TERMINAL-TEST}(s)$
- Now: $\text{H-MINIMAX}(s) = \text{EVAL}(s)$ if $\text{CUTOFF-TEST}(s)$
- i.e. run minimax until depth d , then use evaluation function to choose nodes
- Can also consider iterative deepening approach

Stochastic Games

- How to deal with games with *randomisation*?
- Game tree now has added *chance layers* — even more complex
- Calculating the expected value of a state — MUCH harder than deterministic games

5 Lecture 5: Constraint Satisfaction Problems

AIMA Chapter 6.1-6.4

5.1 CSP Formulation

CSP representation

- Variables $\vec{X} = X_1, \dots, X_n$
- Domain D for variables — X_i has domain D_i — list of values a variable can take
- Constraints \vec{C} — restrictions on values a variable can take
 - Defined by constraint language: algebra/logic (don't give abstract english description)

CSP objective

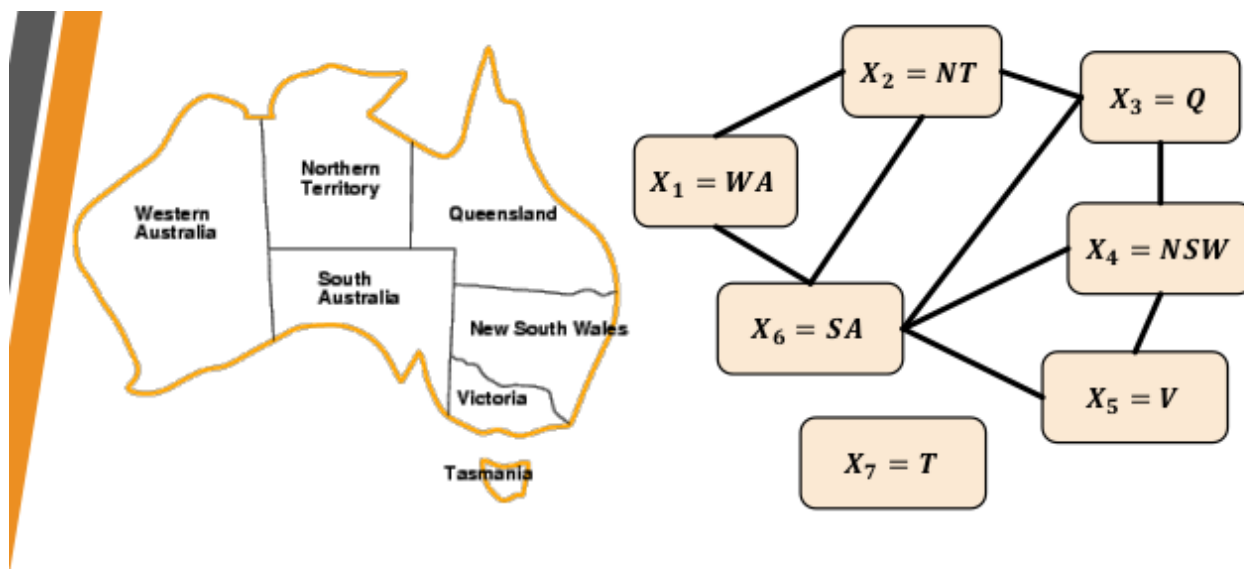
- Find a legal assignment $(y_1, \dots, y_n) — y_i \in D_i$ for all $i \in n$
- *Complete*: all variables assigned values
- *Consistent*: all constraints satisfied

5.1.1 Example: Graph Colouring

Constraint graph: node are variables, edges are constraints

- Variables: $\vec{X} = \langle WA, NT, Q, NSW, V, SA, T \rangle$
- Domains: $D_i = \{R, G, B\}$
- Constraints: if $(X_i, X_j) \in E$ then $color(X_i) \neq color(X_j)$

Binary constraint: involves 2 variables



Variables:	$\vec{X} = \langle WA, NT, Q, NSW, V, SA, T \rangle$
Domains:	$D_i = \{R, G, B\}$
Constraints:	If $(X_i, X_j) \in E$ then $color(X_i) \neq color(X_j)$

5.1.2 Example: Job-Shop Scheduling

- Car assembly consists of 15 tasks
- Variables: $Axle_F$, $Axle_B$, $Wheel_{LF}$, $Wheel_{RF}$, $Wheel_{LB}$, $Wheel_{RB}$, $Nuts_{LF}$, $Nuts_{RF}$, $Nuts_{LB}$, $Nuts_{RB}$, Cap_{LF} , Cap_{RF} , Cap_{LB} , Cap_{RB} , $Inspect$
- Domain: $D_i = \{1, 2, \dots, 27\}$
- Precedence constraints: e.g. $Axle_F + 10 \leq Wheel_{RF}$
- Disjunctive constraints: e.g. $(Axle_F + 10 \leq Axle_B) \text{ or } (Axle_B + 10 \leq Axle_F)$

5.2 CSP Variants

Discrete variables

- Finite domains: e.g. sudoku
- Infinite domains: integers, strings etc. e.g. job-shop scheduling

Continuous variables

- E.g. start/end times for Hubble Space Telescope observations
- Linear programming problems can be solved in polynomial time


5.3 Constraint Variants

- Unary constraints: 1 variable e.g. $SA \neq Green$
- Binary constraints: 2 variables e.g. $SA \neq WA$
- Global/higher-order constraints: 3 or more variables e.g. $X_1 + X_2 - 4X_7 \leq 15$

5.3.1 Example: Cryptarithmic Puzzle

- Each letter represents one digit (base 10)
- Different letters should correspond to different digits
- T and F cannot be 0

$$\begin{array}{r} TWO \\ + TWO \\ \hline FO UR \end{array}$$



Variables:	$\vec{X} = \langle F, T, U, W, R, O, C_1, C_2, C_3 \rangle$
Domains:	$D_i = \{0, \dots, 9\}$
Constraints:	$AllDiff(F, T, U, W, R, O)$ $O + O = R + 10C_1$ $C_1 + W + W = U + 10C_2$ $C_2 + T + T = O + 10C_3$ $C_3 = F$ $T, F \neq 0$

(Also, C_1, C_2, C_3 should be either 0 or 1)

Drawing constraints

- Global constraints: draw using square
 - E.g. $AllDiff(F, T, U, W, R, O)$ — one square linking them all
- Binary constraints: can draw using square, if not just draw an edge directly
- Unary constraints: don't need to draw

5.3.2 Example: Sudoku

	1	2	3	4	5	6	7	8	9
A			3		2		6		
B	9			3		5			1
C			1	8		6	4		
D			8	1		2	9		
E	7								8
F			6	7		8	2		
G			2	6		9	5		
H	8			2		3			9
I			5		1		3		

(a)

	1	2	3	4	5	6	7	8	9
A	4	8	3	9	2	1	6	5	7
B	9	6	7	3	4	5	8	2	1
C	2	5	1	8	7	6	4	9	3
D	5	4	8	1	3	2	9	7	6
E	7	2	9	5	6	4	1	3	8
F	1	3	6	7	9	8	2	4	5
G	3	7	2	6	8	9	5	1	4
H	8	1	4	2	5	3	7	6	9
I	6	9	5	4	1	7	3	8	2

(b)

Variables:	$A_1, \dots, A_9, \dots, I_1, \dots, I_9$ (81 variables)
Domains:	$D_i = \{1, \dots, 9\}$
Constraints:	$AllDiff(\dots) \times 27$ (9 columns, 9 rows, 9 boxes) e.g. $AllDiff(A_1, \dots, A_3; B_1, \dots, B_3; C_1, \dots, C_3)$ is the constraint for the top-left box.

15

5.4 CSP Search

5.4.1 Search Formulation

- State and initial state: initially empty assignment \square
- Transition function: assign a valid value to an unassigned variable, fail if no valid assignments
- Goal test: all variables assigned
- Every solution appears at exactly depth n
- Search path is irrelevant

5.4.2 Search Tree

Each level: pick any remaining variable, give it any possible assignment.

Maximum size i.e. total number of leaves: $n! \times d^n$

- E.g. 4 Variables and 3 Values — size = $4! \times 3^4 = 1944$

5.5 Backtracking Search Algorithm

Backtracking search

- More efficient than the search above
- Perform DFS with single-variable/level assignments: at every level, consider assignments to a *single* variable

- Order of variable assignment is irrelevant

BACKTRACKING-SEARCH(*csp*) returns a solution, or failure

- return BACKTRACK($\{\}$, *csp*)


BACKTRACK(*assignment*, *csp*) returns a solution, or failure

- if assignment is complete, return it
- $\text{var} \leftarrow \text{SELECT-UNASSIGNED-VARIABLE}(csp)$
- for each value in ORDER-DOMAIN-VALUES(*var*, *assignment*, *csp*):
 - if value is consistent with assignment:
 - * add $\{var = value\}$ to assignment
 - * inferences $\leftarrow \text{INFERENCE}(csp, var, value)$
 - * if inferences == failure: continue
 - * add inferences to assignment
 - * result $\leftarrow \text{BACKTRACK}(assignment, csp)$
 - * if result \neq failure: return result
 - remove $\{var = value\}$ and inferences from assignment
- return failure

5.6 Backtracking Heuristics: Variable and Value Ordering

5.6.1 Variable-Order Heuristics: SELECT-UNASSIGNED-VARIABLE


1. Most constraining variable a.k.a. degree heuristic: choose the variable that imposes the most constraints on the remaining unassigned variables
 - This is best: it reduces the branching factor \Rightarrow likely get to terminal state faster



1. Most constrained variable a.k.a. Minimum-Remaining-Values (MRV) heuristic: choose the variable with the fewest remaining legal values
 - Use as tiebreaker

5.6.2 Value-Order Heuristic: ORDER-DOMAIN-VALUES

1. Least constraining value: choose the value that rules out the fewest values for the neighbouring unassigned variables
 - Because we're "actually trying to solve the problem" in this stage, unlike the variable stage



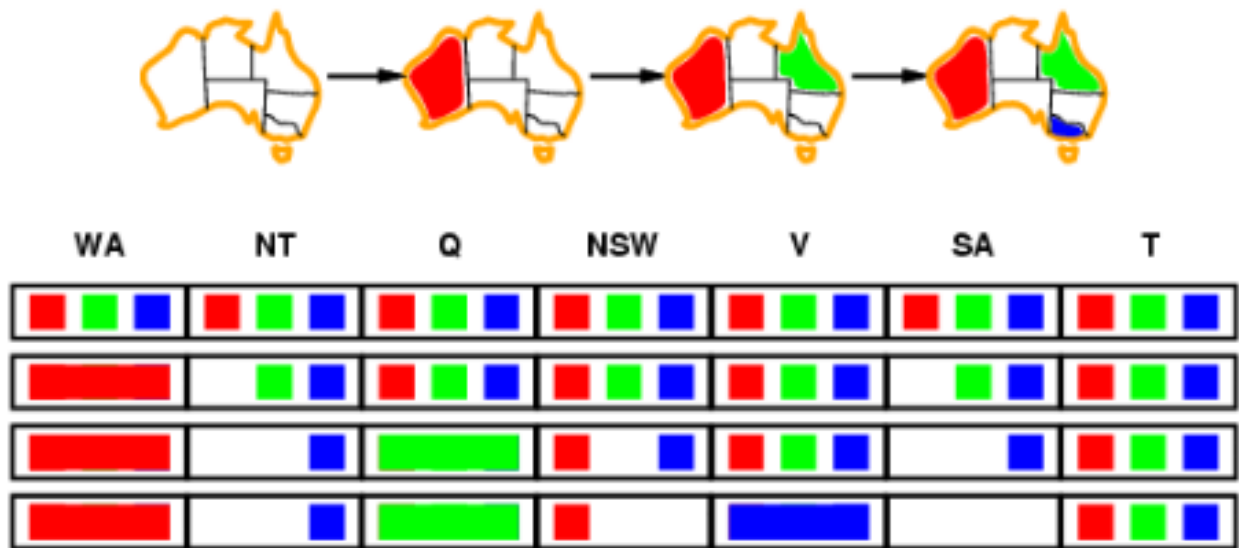
5.7 Inference: INFERENCE

Idea: check for failures early.

5.7.1 Forward Checking

Forward checking

- Keep track of remaining legal values for unassigned variables
- (★) Terminate search when any variable has no legal values left




E.g. here SA has no remaining valid assignments => failure.

5.7.2 Constraint Propagation

Constraint propagation: 'move ahead' with the constraints

- Repeatedly locally enforce constraints
- Infer illegal values for assignments early on



E.g. here NT and SA both have to be blue, but by constraints, they can't be both blue

6 Lecture 6: Project Details

6.1 Reinforcement Learning

1. \leftarrow Agent receives input information
2. \rightarrow Agent performs valid action
3. \leftarrow Agent obtains reward

State $s_t \rightarrow$ action $a_t \rightarrow$ reward r_t (also takes you to state s_{t+1})

6.2 Supervised vs Unsupervised Learning

Unsupervised: data are unlabelled \Rightarrow perform things like clustering

Supervised: data are labelled \Rightarrow perform things like predicting labels for new unlabelled data

- Classification problems: supervised learning problem with discrete-valued class

#	x_1	\dots	x_q	y
1	$x_{1,1}$	\dots	$x_{1,q}$	y_1
\vdots	\vdots		\vdots	\vdots
p	$x_{p,1}$	\dots	$x_{p,q}$	y_p

Goal: build a model F such that $F(X) = y$ with high accuracy, where X is a new unseen instance

6.3 Evaluating Classification Models

Generating models and evaluating models are different!

Idea behind evaluation: measure generalisation performance

- Assume instances are governed by overarching distribution D
- Want to determine, the probability of accurately classifying ANY instance drawn from D

Example methods

- Hold-out testing, i.e. training and testing sets
- k-fold cross-validation

6.4 Algorithm Selection

Given a classification dataset S , and a set of algorithms we'll use A , determine which algorithm $a^* \in A$ is optimal

Meta-learning: pose algorithm selection problem as another classification problem

Generate meta-dataset

- Each x_i corresponds to a *characteristic* of a dataset (e.g. number of instances, r^2 , mutual information, etc.)
- Each y_i corresponds to the optimal algorithm a^* for that dataset

Problem: what is the *overarching distribution* governing the algorithm selection problem?


- Which datasets are properly representative for this problem?
- Where can we draw them from?
- Repositories exist, but are these representative of all possible problems?

Just ensure that the model built has *good coverage*; uniform distribution of datasets

- At least have many instances representing different patterns of when one algorithm will be better than another

Project 1 Search Problem

Search for some S^* such that when we plot each $s_j \in S^*$ within the expertise space, we get a distribution that is close to uniform



7 Lecture 7: Constraint Satisfaction Problems II


7.1 Inference in CSPs: Arc Consistency and AC-3

Constraint propagation: node consistency for unary constraints, arc consistency for binary constraints


7.1.1 Arc Consistency

Arc consistency: X is arc-consistent wrt X_j i.e. arc (X_i, X_j) is consistent, iff for every value $x \in D_i$ there exists some value $y \in D_j$ that satisfies binary constraint on arc (X_i, X_j)

- Note that arcs are *directed*.
- To maintain AC: remove a value if it makes a constraint impossible to satisfy.



(SA, NSW): OK



(NSW, SA): Need to remove blue value from NSW

After an update on X_i where it loses a value, we MUST *re-check* the neighbours of X_i .



(V, NSW): Now that NSW cannot be blue, V cannot be red

7.1.2 AC-3 Algorithm

AC-3(csp) returns *false* if inconsistency is found, otherwise *true*

- $queue \leftarrow$ all the arcs in csp
- while $queue$ is not empty:
 - $(X_i, X_j) \leftarrow \text{REMOVE-FIRST}(queue)$
 - if REVISE(csp, X_i, X_j):
 - * if size of $D_i = 0$ then return *false*

- * for each X_k in $\text{NEIGHBOURS}(X_i) - \{X_j\}$:

- add (X_k, X_i) to *queue*

REVISE(csp, X_i, X_j) returns *true* if we revise the domain of X_i

- $revised \leftarrow false$
- for each x in D_i :
 - if no value y from D_j allows (x, y) to satisfy constraint between X_i and X_j :
 - * delete x from D_i
 - * $revised \leftarrow true$
- return *revised*

Time complexity: $O(n^2 d^3)$

- CSP has at most n^2 directed arcs
- Each arc (X_i, X_j) can be inserted at most d times into the queue, since X_i has at most d values
- REVISE: checking consistency of arc takes $O(d^2)$ time
- AC-3: $O(n^2 \times d \times d^2) = O(n^2 d^3)$

7.1.3 Maintaining AC (MAC)

Search procedure

- Establish AC at root
- When AC-3 terminates, choose a new variable and value
- Re-establish AC given the new variable choice — maintain AC
- Repeat;
- Backtrack if AC gives *empty domain*

We could use AC-3 purely as preprocessing, or do it at every step

AC-3 with preprocessing

- Add all arcs

AC-3 with backtracking

- If domain of variable X' is updated, then only need to add all arcs leading to X'
- i.e. check each arc (X_i, X')

7.1.4 Generalised Arc Consistency (not covered in CS3243)

What if our arcs are global and not binary constraints?

- Can reduce to binary constraints if possible
- Otherwise, we can extend arc consistency (2-consistency) to k-consistency

8 Logical Agents

AIMA Chapter 7

8.1 Knowledge-based Agents

Previously: we use search; no real model of what the agent knows Now: we represent agent domain knowledge using logical formulas

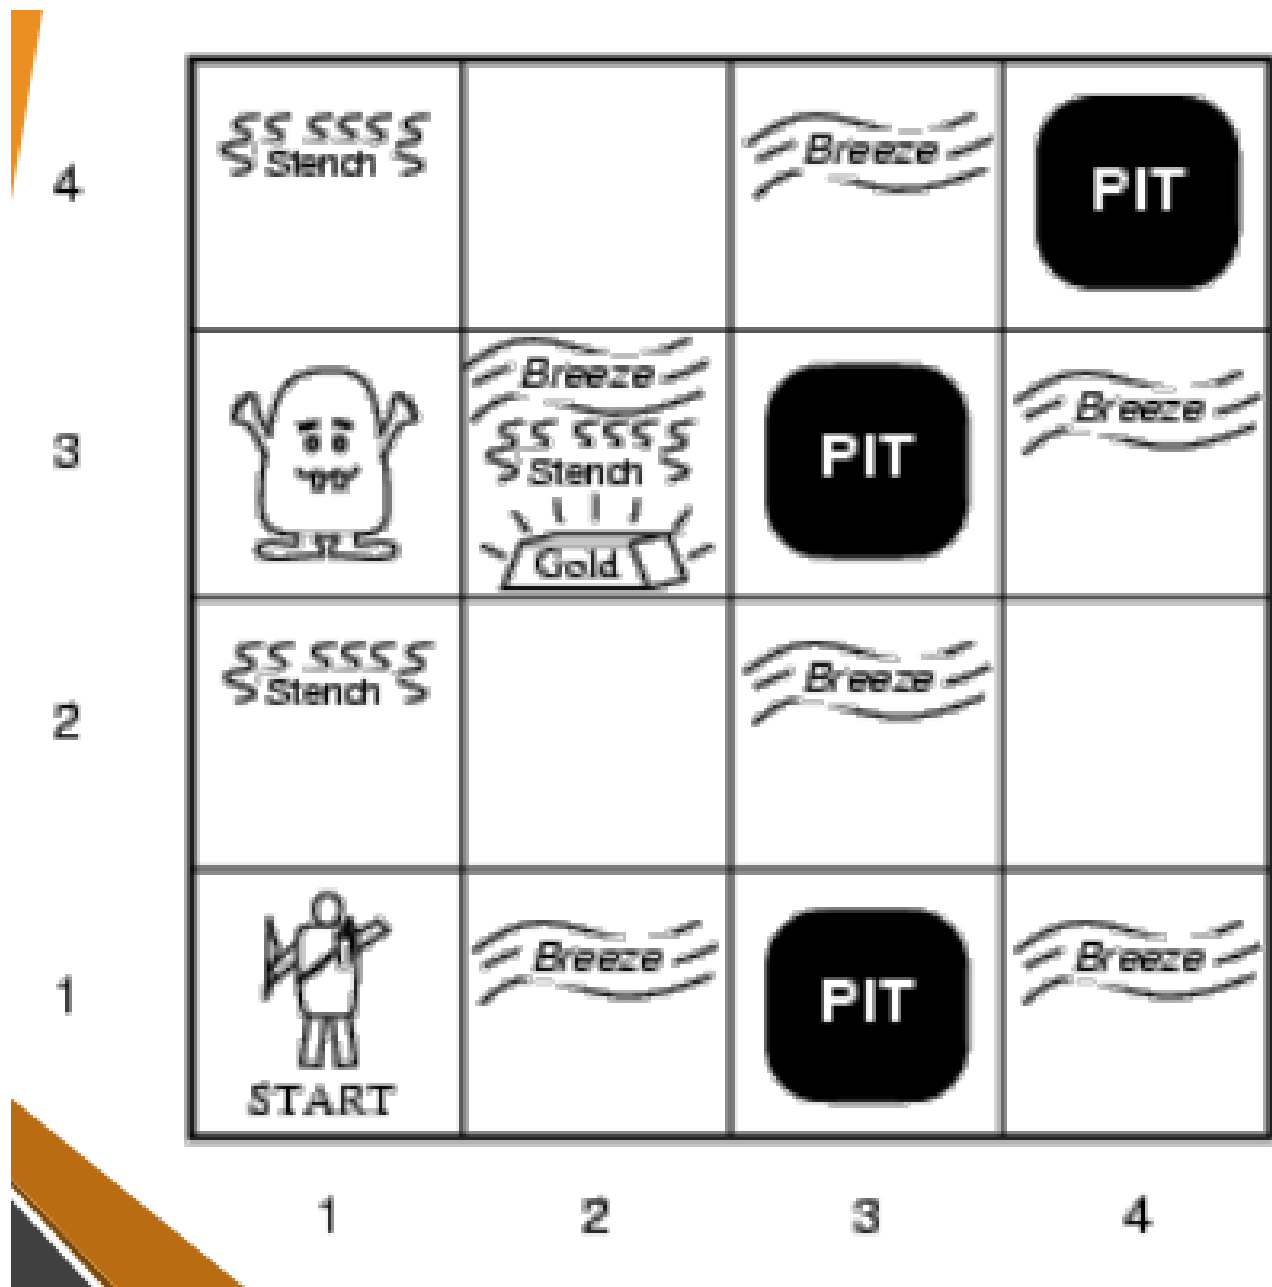
Logical agent: Inference Engine + Knowledge Base

- Inference Engine: domain-independent algorithms
- Knowledge Base: domain-specific content — *set of sentences* in a formal language
 - Pre-populate with background/domain knowledge (e.g. game rules)

KB-AGENT(*percept*) returns an *action*

- persistent: KB , a knowledge base; t , a counter for time initially set to 0
- TELL(KB , MAKE-PERCEPT-SEQUENCE(*percept*, t))
- $action \leftarrow$ ASK(KB , MAKE-ACTION-QUERY(t))
- TELL(KB , MAKE-ACTION-SEQUENCE(*action*, t))
- $t \leftarrow t + 1$
- return *action*

8.2 Example: Wumpus World



Wumpus and pits will kill you

- Beside wumpus: stench
- Beside pit: breeze

Task environment (PEAS)

- Performance measure: +1000 for gold, -1000 for dying, -1 for each action, -10 for using arrow
- Environment: 4x4 grid of rooms
- Actuators: forward, turn left, turn right, grab gold, shoot arrow
- Sensors: perceive stench/breeze/glitter/scream

Environment

- Fully observable: no — only local perception

- Deterministic: yes
- Episodic: no — sequential actions
- Static: yes
- Discrete: yes
- Single-Agent: yes

Initial KB

- If there is a PIT, there is a BREEZE beside it
- If there is a WUMPUS, there is a STENCH beside it
- It's a 4x4 grid world
- ...

8.3 Logic

Logic: formal language for KR, consists of syntax + semantics

- Syntax: defines valid sentences in language: S_1, S_2 , etc.
 - Provides logical connectives for constructing complex sentences from simpler ones, e.g. $S_1 \wedge S_2$ etc.
 - e.g. $x + y = 4$ is a sentence
- Semantics: defines the meaning of a sentence; the "truth of each sentence with respect to each possible world"
 - i.e. defines truth (validity) of a sentence in a given world (for some given value assignments in an environment)
 - e.g. $x + y = 4$ is true in a world where $x = 2$ and $y = 2$, but false in a world where $x = 1$ and $y = 1$

8.3.1 Logical Reasoning: Entailment

Modelling: m models/satisfies sentence α if α is true under m

- A model represents the idea of a "possible world" — it assigns a truth value to all the variables
- Let $M(\alpha)$ be the set of all models satisfying α
- E.g. $\alpha = (q \in \mathbb{Z}_+) \wedge (\forall n, m \in \mathbb{Z}_+ : q = nm \Rightarrow n \vee m = 1)$


Entailment \models : means that one sentence follows logically from another sentence

- $\alpha \models \beta$ is equivalent to $M(\alpha) \subseteq M(\beta)$
- E.g. $\alpha = (q \text{ is prime})$ entails $\beta = (q \text{ is odd}) \vee (q = 2)$

KB is true \Leftrightarrow all its rules are true, i.e. $\bigwedge_{k=1}^n R_k$ is true


Key takeaway: if our model is a subset of a sentence α , then α is true

also true




Example: Wumpus World

- Suppose we move right to (2,1) to detect a breeze
- Consider 8 possible models for KB with pits (3 boolean choices \Rightarrow 8 possible models)
- KB is only true



- Suppose we want to infer sentence $\alpha_1 = "(1,2) \text{ is safe}"$.
- True: proved by model checking. Worlds satisfying KB \subseteq worlds where (1,2) is safe



- Let P_{ij} be whether there's a pit in (i, j) .
- Let B_{ij} be whether there's a breeze in (i, j) .

Rules

- $R_1 : \neg P_{1,1}$
- $R_4 : \neg B_{1,1}$
- $R_5 : P_{2,1}$

"Pits cause breezes in adjacent squares"

- $R_2 : B_{1,1} \Rightarrow (P_{1,2} \vee P_{2,1})$
- $R_3 : B_{2,1} \Rightarrow (P_{1,1} \vee P_{2,2} \vee P_{3,1})$

8.3.2 Inference Algorithm

Let $KB \vdash_A \alpha$ be "sentence α is derived/inferred from KB by inference algorithm A ".

- A is *sound* if $KB \vdash_A \alpha$ implies $KB \models \alpha$
 - If KB derives α , then KB entails α
 - Whatever is derived is correct, i.e. "don't infer nonsense"
- A is *complete* if $KB \models \alpha$ implies $KB \vdash_A \alpha$
 - If KB entails α , then KB derives α
 - Whatever is correct is derived, i.e. if it's implied it will be inferred

We want an inference algorithm that is both *sound* and *complete*.

- Let X = all sentences derived from KB using A
- Let Y = all possible sentences entailed by KB
- $X = Y$: sound and complete
- $X \subset Y$: sound and not complete
- $Y \subset X$: not sound and complete
- Otherwise: not sound and not incomplete

8.3.3 Inference!

- Given a knowledge base, infer something about the world
- Inference: deriving new knowledge out of percepts
- Given KB and α , we want to know if $KB \models \alpha$, i.e. can we infer α from KB ?

8.3.4 Truth Table for Inference

Truth Table for Inference

$B_{1,1}$	$B_{2,1}$	$P_{1,1}$	$P_{1,2}$	$P_{2,1}$	$P_{2,2}$	$P_{3,1}$	KB	α_1	
false	false	false	false	false	false	false	false	true	
false	false	false	false	false	false	true	false	true	
⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	
false	true	false	false	false	false	false	false	true	
false	true	false	false	false	false	true	true	true	KB true
false	true	false	false	false	true	false	true	true	
false	true	false	false	false	true	true	true	true	
false	true	false	false	true	false	false	false	true	
⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	
true	true	true	true	true	true	true	false	false	

- Build a truth table of all possible values
- Evaluate the models where the KB is true
- Does KB entail α_1 : See if the remaining rows are true for α_1 . If so, we can infer it

Inference by Truth Table Enumeration

- Sound: directly implements entailment, and calculates all possible inferences from KB by brute force
- Complete: only finitely many combinations of truth assignments, and goes through all
- (For above 2, see diagnostic quiz 8/Sam's slides W10)
- Time complexity: $O(2^n)$

- Space complexity: $O(n)$ — because the enumeration is depth-first

8.4 Validity and Satisfiability

Validity: a sentence is *valid* if it is true in *all* models

- E.g. statements that are true regardless of truth assignments (tautology), e.g. *True*, $A \vee \neg A$
- $KB \models \alpha$ iff $(KB \Rightarrow \alpha)$ is valid

Satisfiability: a sentence is *satisfiable* if it is true in *some* model

Unsatisfiability: a sentence is *unsatisfiable* if it is true in *no* models

- $KB \models \alpha$ iff $(KB \wedge \neg \alpha)$ is unsatisfiable

8.5 Applying Inference Rules

Form of search problem: search for more knowledge (search grows our KB)

- States: KBs
- Actions: inference rules
- Transition: add sentence to current KB
- Goal: KB contains sentence to prove

Examples of inference rules

- And-elimination: $a \wedge b \models a$
- Modus ponens: $a \wedge (a \Rightarrow b) \models b$
- Logical equivalences: $(a \vee b) \models \neg(\neg a \wedge \neg b)$

8.6 Resolution (for CNF)

CNF: conjunction of disjunctions i.e. 'and's of 'or's

- E.g. $(x_1 \vee \neg x_2) \wedge (x_2 \vee x_3 \vee \neg x_4)$
- Conversion to CNF: simple standard stuff

Resolution: if x appears in C_1 and $\neg x$ appears in C_2 , it can be deleted (x must be a literal)

- $(P \vee x) \wedge (Q \vee \neg x)$ is the same as $(P \vee Q)$
- Resolution is *sound* and *complete* for propositional logic

(★) Resolution algorithm

- Proof by contradiction: to prove α , suppose otherwise add $\neg \alpha$ into the KB
- Step 1: add $\neg \alpha$ into KB
- Step 2: convert KB to CNF
- Step 3: pick 2 rules and reduce; repeat
- Use resolution to see if the eventual KB is \emptyset i.e. contradiction

Resolution algorithm is sound and complete

- Soundness: why (???)
- Completeness: why (???)

8.6.1 Example

Assume we are at (1,1), and we want to infer if there is no pit at (1,2)

- $KB = \neg B_{1,1} \wedge B_{1,1} \Leftrightarrow (P_{1,2} \vee P_{2,1})$
- $\alpha = \neg P_{1,2}$

Resolution algorithm

- Step 1: add $\neg\alpha$ to KB
 - $KB = \neg B_{1,1} \wedge (B_{1,1} \Leftrightarrow (P_{1,2} \vee P_{2,1})) \wedge P_{1,2}$
- Step 2: convert KB to CNF
 - $KB = \neg B_{1,1} \wedge P_{1,2} \wedge (\neg B_{1,1} \vee P_{1,2} \vee P_{2,1}) \wedge (\neg P_{1,2} \vee B_{1,1}) \wedge (\neg P_{2,1} \vee B_{1,1})$
- Step 3: pick two rules and reduce
 - Reduce rule 2 and rule 4: $P_{1,2}$ in rule 2 and $\neg P_{1,2}$ in rule 4
 - Reduced to rule 6: $B_{1,1}$
 - Reduce rule 1 and rule 6: $\neg B_{1,1}$ in rule 1 and $B_{1,1}$ in rule 6
 - Reduce to \emptyset

8.7 KB and Horn Clauses

Horn clauses: of form $B_1 \wedge B_2 \wedge \dots \wedge B_k \Rightarrow A$

- Forward/backward chaining is *sound* and *complete* for KB comprised of horn clauses

Clauses with at most 1 positive literal

- Clause is a sentence comprising disjunctions: e.g. $A \vee \neg B$, $\neg A \vee \neg C \vee D$

Three forms of horn clauses

- Literals (facts): e.g. A
- Definite clause (rules): e.g. $B_1 \wedge B_2 \wedge \dots \wedge B_k \Rightarrow A$ i.e. $\neg B_1 \vee \neg B_2 \vee \dots \vee \neg B_k \vee A$

8.8 Forward Chaining

Idea: keep adding literals/facts


Idea: fire any rule whose premise is satisfied in the KB, add its conclusion to the KB, repeat until query Q is found

AND-OR graph

KB of horn clauses

$$\begin{aligned} P &\Rightarrow Q \\ L \wedge M &\Rightarrow P \\ B \wedge L &\Rightarrow M \\ A \wedge P &\Rightarrow L \\ A \wedge B &\Rightarrow L \\ A \\ B \end{aligned}$$

AND-OR graph



FC algorithm

- For every rule c , let $count(c)$ be the number of literals in its premise
- For every literal s , let $inferred(s)$ be initially false
- Let $agenda$ be a queue of literals, initially containing all literals known to be true
- While $agenda \neq \emptyset$:
 - Pop literal p from $agenda$; if it is Q , we are done
 - Set $inferred(p)$ to be true
 - For each clause $c \in KB$ such that p is in the premise of c , decrement $count(c)$
 - If $count(c) = 0$, add conclusion of c to $agenda$

Example

- Iteration 1: $agenda = [A, B]$
- Iteration 2: $agenda = [B]$
- Iteration 3: $agenda = [] \Rightarrow [L]$
- Iteration 4: $agenda = [] \Rightarrow [M]$
- Iteration 5: $agenda = [] \Rightarrow [P]$
- Iteration 6: $agenda = [] \Rightarrow [Q]$

Proof of completeness

- FC derives every atomic sentence/literal entailed by a horn KB
- Suppose FC reaches a fixed point, where no new atomic sentences are derived
- Consider the final state as a model m that assigns true/false to symbols based on inferred table
- Every clause in the original KB is true in m
- Hence m is a model of KB

- If $KB \models q$, then q is true in *every* model of KB, including m

8.9 Backward Chaining

Idea: work backwards from the query Q

To prove Q by backwards chaining,

- Check if Q is known already, or
- Prove by backwards chaining the premise of some rule concluding in Q
- We need to avoid loops: check if the new subgoal is already on the goal stack
- Backtracking DFS

8.10 Forward vs Backward Chaining

- FC: data-driven reasoning
 - When you don't know the goal, but want to try to build towards it
 - May do a lot of work that is irrelevant to the goal
- BC: goal-driven reasoning
 - When you know the goal, and want to work backwards to prove it
 - Complexity of BC can be sublinear in size of KB

9 Uncertainty

9.1 Probability Basics

Probability

- Random variable X : quantifies an outcome of a random occurrence
- Domain D_X : set of all outcomes of a random variable
- Event: subset of a domain
- Probability distribution: assigns a probability value $p(x) \in [0, 1]$ to every $x \in D_X$

Axioms of probability

- Total probability is 1: $\sum_{x \in D_X} p(x) = 1$
- $P(A \cup B) = P(A) + P(B) - P(A \cap B)$

Multiple random variables

- Joint probability: $p(x, y) = P(X = x, Y = y)$ (discrete)
- Marginal probability: $p(x) = \sum_{y \in D_Y} p(x, y)$
- Conditional probability: e.g. $P(A|B) = \frac{P(A \cap B)}{P(B)}$

Bayes' rule: $P(A|B) = \frac{P(B|A) \times P(A)}{P(B)}$ Chain rule: $P(X_1, X_2, \dots, X_k) = \prod_{i=1}^k P(X_i | X_1, \dots, X_{i-1})$

Independence

- $P(A \cap B) = P(A) \times P(B)$, i.e. $P(A|B) = P(A)$
- Conditional independence: $P(X \cap Y | Z) = P(X | Z) \times P(Y | Z)$

9.2 Bayesian Inference

$P(X | Y_1, \dots, Y_k)$ — we want to find the probability of event X , given probabilities of other events Y_i

Inference by enumeration

- Find $P(X)$ by summing over all atomic events
- $P(X) = \sum_{x \in X} P(X = x)$

Bayes' rule and conditional independence

- $P(C | E_1, \dots, E_n) = \frac{P(C) \times P(E_1, \dots, E_n | C)}{P(E_1, \dots, E_n)} \propto \prod_{i=1}^n P(E_i | C)$
- This is an example of the naive Bayes' model

Normalisation

- $P(X | Y_1, Y_2) = \frac{P(Y_1, Y_2 | X) \times P(X)}{P(Y_1, Y_2)}$
- But we don't care about $P(Y_1, Y_2)$, so set it to α
- Then $P(X = \text{healthy} | A) = \alpha \times P(X = \text{healthy}) \times P(Y_1 = y_1 | X = \text{healthy}) \times P(Y_2 = y_2 | X = \text{healthy}) = \dots$
- Then $P(X = \text{sick} | A) = \alpha \times P(X = \text{sick}) \times P(Y_1 = y_1 | X = \text{sick}) \times P(Y_2 = y_2 | X = \text{sick}) = \dots$

10 Bayesian Networks

Represent joint distributions via a graph

- Nodes: random variables
- Edges: assume X causes/influences Y
- For each node X , we can get a conditional distribution for X given its parents, i.e. $P(X|Parents(X))$
- Conditional probability table (CPT): the conditional distribution of X for each combination of parent values

Then $P(X_1, \dots, X_n) = \prod_{i=1}^n P(X_i|Parents(X_i))$

- The fewer parents overall, the better (the less complex the graph is)

Complexity


- If each variable has $\leq k$ parents, then network representation requires $O(n2^k)$ values, compared to $O(2^n)$ for full joint distribution

10.1 Examples

Example: independent causes/common effect


$$P(A, B, C) = P(C|A, B) \cdot P(A) \cdot P(B)$$

- A and B are pairwise independent, *unless* you condition on observing the effect C : then A and B are conditionally dependent



Example: independent events


$$P(A, B, C) = P(A) \cdot P(B) \cdot P(C)$$



Example: conditionally independent effects/common cause

$$P(A, B, C) = P(C|A) \cdot P(B|A) \cdot P(A)$$


- B and C are conditionally independent given A



Example: causal chain

$$P(A, B, C) = P(C|B) \cdot P(B|A) \cdot P(A)$$

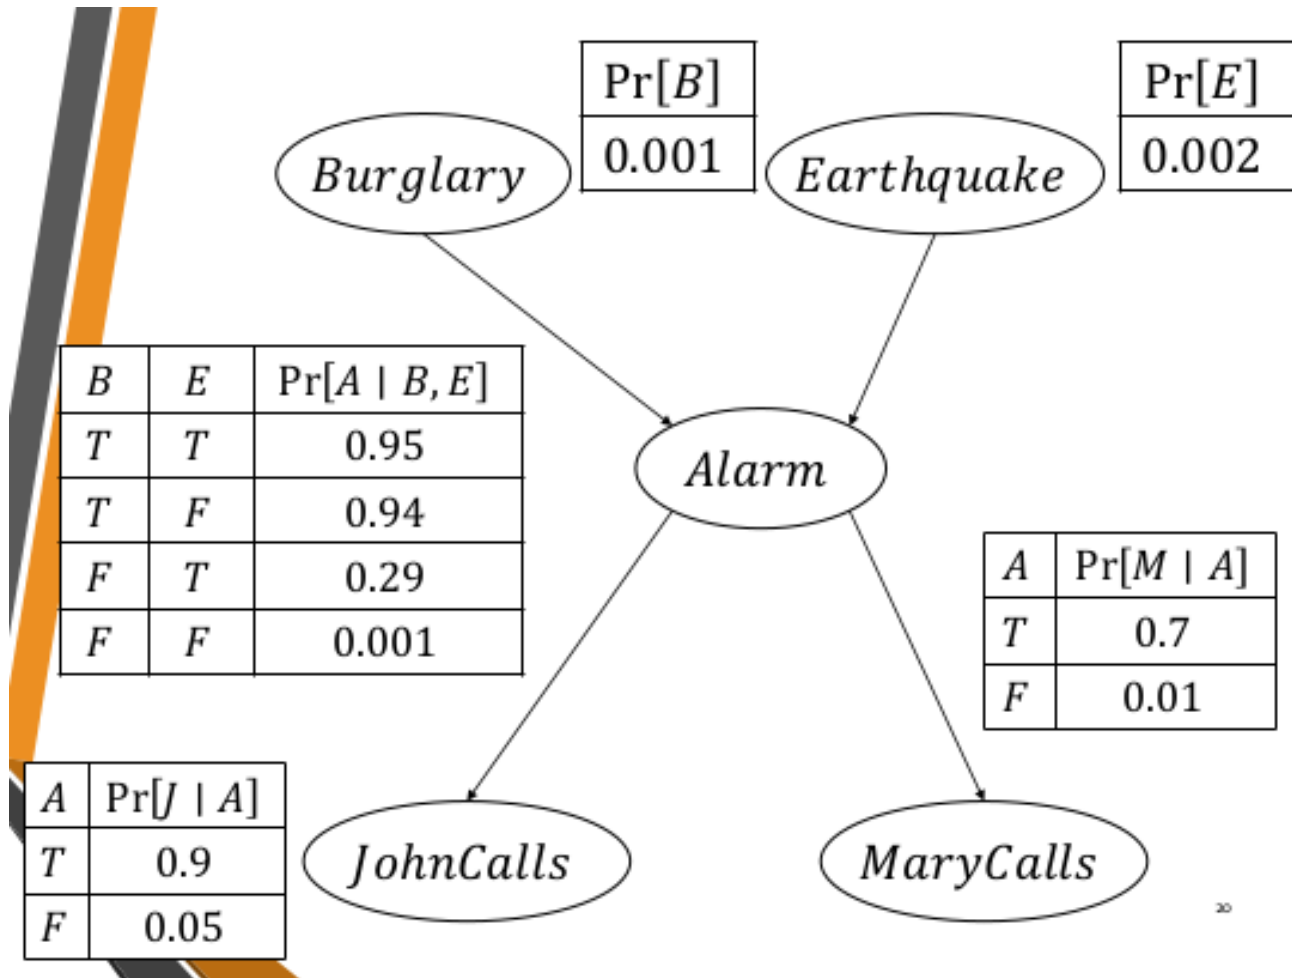
- C is conditionally independent of A given B – note that $P(C|B) = P(C|B, A)$



Example: burglary

- A : Alarm goes off
- E : Alarm sometimes set off by minor earthquake
- B : Alarm set off by burglar

- J : John calls to say my house alarm is ringing
- M : Mary calls to say my house alarm is ringing



$$P(B = 1 | J = 1, M = 0) = \frac{P(B=1, J=1, M=0)}{P(J=1, M=0)} = ?$$

- To find $P(B = 1, J = 1, M = 0)$: sum over 4 cases of $A=0/1, E=0/1$
- To find $P(J = 1, M = 0)$: sum over 8 cases of $A=0/1, E=0/1, B=0/1$
- whereby $P(J, M, A, B, E) = P(J|A) \cdot P(M|A) \cdot P(A|B, E) \cdot P(B) \cdot P(E)$

10.2 Inference in Bayesian Networks

Bayesian network represents the full joint distribution.

Infer any query by summing over all cases of the other variables.

10.3 Algorithm for Constructing Bayesian Network

Algorithm

- Choose an ordering for variables X_1, \dots, X_n
- For $i=1$ to n :
 - Add node X_i to network
 - Select minimal set of parents from X_1, \dots, X_{i-1} such that $P(X_i | Parents(X_i)) = P(X_i | X_1, \dots, X_{i-1})$
 - Add edges from every parent to X_i

- Write down CPT for $P(X_i | \text{Parents}(X_i))$


Variable order matters!

- Choosing a 'good' variable order can reduce the number of edges required

10.4 Markov Blanket

A node is conditionally independent of everything else given the values of its:

- Parents
- Children
- Children's parents



10.5 d-Separation

Given variables X and Y and known variables $\epsilon = \{E_1, \dots, E_k\}$, are X and Y surely independent given ϵ ?

Idea: any general graph can be broken down into three cases (causal chain/common cause/common effect) to determine conditional independence of X and Y given knowledge of ϵ

Check every undirected path between X and Y , ignoring direction of arcs






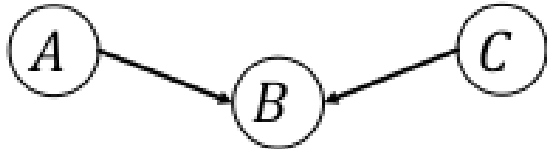
- (\star) If all paths are not active, then X and Y are independent given ϵ

Active path: Path is active iff every triple on the path is active

- I.e. if *any* triple on the path is inactive, the *entire path* is inactive

Active triple: see the chart

- Dark means we know B , light means we don't know B
- Note: only take into account knowledge of B , not A or C in these triples

Active	Inactive
	
	
	

Example

- Here, all 3 potential paths are inactive => 2 red-marked nodes are independent given ϵ

